

# Computational Intelligence

## Definition Sheet

Thomas Gabor                      Maximilian Zorn  
LMU Munich                      LMU Munich

Claudia Linnhoff-Popien  
LMU Munich

winter term 2024/2025

**Notation.**  $\mathbb{B}$  is the set of truth values or Booleans, i.e.,  $\mathbb{B} = \{0, 1\}$ .  $\mathbb{N}$  is the set of natural numbers starting from zero, i.e.,  $\mathbb{N} = \{0, 1, 2, \dots\}$  so that it holds that  $\mathbb{B} \subset \mathbb{N} \subset \mathbb{Z} \subset \mathbb{R} \subset \mathbb{C}$ .  $\mathbb{P}$  denotes the space of probabilities and  $\mathbb{F}$  denotes the space of fuzzy values with  $\mathbb{P} = \mathbb{F} = [0; 1] \subset \mathbb{R}$  being only discerned for semantic but not mathematical reasons.  $\mathbb{E}$  denotes the expected value.

$\wp(X) = 2^X$  denotes the power set of  $X$ .  $\#$  denotes vector or sequence concatenation, i.e., given two vectors  $\mathbf{x} = \langle x_1, \dots, x_{|\mathbf{x}|} \rangle$  and  $\mathbf{y} = \langle y_1, \dots, y_{|\mathbf{y}|} \rangle$ ,  $\mathbf{x} \# \mathbf{y} = \langle x_1, \dots, x_{|\mathbf{x}|}, y_1, \dots, y_{|\mathbf{y}|} \rangle$ . A vector  $\langle x_0, \dots, x_{n-1} \rangle$  with length  $n \in \mathbb{N}$  can also be written as  $\langle x_i \rangle_{0 \leq i \leq n-1}$  for a new iteration variable  $i$ . A sequence containing an arbitrary amount of elements from the space  $X$  has the type  $\langle X \rangle$ .  $_$  denotes unspecified function arguments ( $f(\_) = 0$  is the constant function that always returns zero, e.g.). For any finite set  $X = \{x_0, \dots, x_n\}$ ,  $|X| = n$  denotes the number of elements in  $X$ . For infinite sets,  $|X| = \infty$ .

**Definition 1** (agent). Let  $\mathcal{A}$  be a set of actions. Let  $\mathcal{O}$  be a set of observations. An agent  $A$  can be given via a policy function  $\pi : \mathcal{O} \rightarrow \mathcal{A}$ . Given a time series of observations  $\langle o_t \rangle_{t \in \mathcal{Z}}$  for some time space  $\mathcal{Z}$  the agent can thus generate a time series of actions  $\langle a_t \rangle_{t \in \mathcal{Z}}$  by applying  $a_t = \pi(o_t)$ .

**Algorithm 1** (brute force (policy)). Let  $\mathcal{A}$  be a set of actions. Let  $\mathcal{O}$  be a set of observations. Let  $\Gamma \subseteq (\mathcal{O} \rightarrow \mathcal{A}) \rightarrow \mathbb{B}$  be a space of goal predicates on policy functions. Let  $\gamma \in \Gamma$  be a goal predicate. We assume that the policy space  $\Pi \subseteq \mathcal{O} \rightarrow \mathcal{A}$  is enumerable, i.e.,  $\Pi = \langle \pi_i \rangle_{i \in \mathbb{N}}$ . Brute force starting from  $i$  is given via the function

$$b(i) = \begin{cases} \pi_i & \text{if } \gamma(\pi_i), \\ b(i+1) & \text{otherwise.} \end{cases}$$

If not further specified, the call to  $b(0)$  is called brute force search for an agent policy. Usually, an additional termination condition is specified.

**Algorithm 2** (random search (policy)). Let  $\mathcal{A}$  be a set of actions. Let  $\mathcal{O}$  be a set of observations. Let  $\Gamma \subseteq (\mathcal{O} \rightarrow \mathcal{A}) \rightarrow \mathbb{B}$  be a space of goal predicates on policy functions. Let  $\gamma \in \Gamma$  be a goal predicate. We assume that the policy space  $\Pi \subseteq \mathcal{O} \rightarrow \mathcal{A}$  can be sampled from, i.e.,  $\pi \sim \Pi$  returns a random element from  $\Pi$ . Random search for  $n$  samples is given via the function

$$\rho(n) = \begin{cases} \emptyset & \text{if } n = 0, \\ \pi & \text{if } n > 0 \text{ and } \gamma(\pi) \text{ where } \pi \sim \Pi, \\ \rho(n-1) & \text{otherwise.} \end{cases}$$

**Definition 2** (optimization). Let  $\mathcal{X}$  be an arbitrary state space. Let  $\mathcal{T}$  be an arbitrary set called target space and let  $\leq$  be a total order on  $\mathcal{T}$ . A total function  $\tau : \mathcal{X} \rightarrow \mathcal{T}$  is called target function. Optimization (minimization/maximization) is the procedure of searching for an  $x \in \mathcal{X}$  so that  $\tau(x)$  is optimal (minimal/maximal). Unless stated otherwise, we assume minimization. An optimization run of length  $g+1$  is a sequence of states  $\langle x_t \rangle_{0 \leq t \leq g}$  with  $x_t \in \mathcal{X}$  for all  $t$ .

Let  $e : \langle \mathcal{X} \rangle \times (\mathcal{X} \rightarrow \mathcal{T}) \rightarrow \mathcal{X}$  be a possibly randomized or non-deterministic function so that the optimization run  $\langle x_t \rangle_{0 \leq t \leq g}$  is produced by calling  $e$  repeatedly, i.e.,  $x_{t+1} = e(\langle x_u \rangle_{0 \leq u \leq t}, \tau)$  for all  $t$ ,  $1 \leq t \leq g$ , where  $x_0$  is given externally (e.g.,  $x_0 =_{def} 42$ ) or chosen randomly (e.g.,  $x_0 \sim \mathcal{X}$ ). An optimization process is a tuple  $(\mathcal{X}, \mathcal{T}, \tau, e, \langle x_t \rangle_{0 \leq t \leq g})$ .

**Definition 3** (optimization (policy)). Let  $\mathcal{X} = \Pi$  be a policy space. Let  $\mathcal{D} = (\Pi, \mathcal{T}, \tau, e, \langle x_t \rangle_{0 \leq t \leq g})$  be an optimization process according to Definition 2.  $\mathcal{D}$  is called a policy optimization process.

**Algorithm 3** (simulated annealing). Let  $\mathcal{D} = (\mathcal{X}, \mathcal{T}, \tau, e, \langle x_u \rangle_{0 \leq u \leq t})$  be an optimization process. Let  $neighbors : \mathcal{X} \rightarrow \wp(\mathcal{X})$  be a function that returns a set of neighbors of a given state  $x \in \mathcal{X}$ . Let  $\kappa : \mathbb{N} \rightarrow \mathbb{R}$  be a temperature schedule, i.e., a function that returns a temperature value for each time step. Let  $A : \mathcal{T} \times \mathcal{T} \times \mathbb{R} \rightarrow \mathbb{P}$  with  $\mathbb{P} = [0; 1] \subset \mathbb{R}$  be a function that returns an acceptance probability given two target values and a temperature. Commonly, we use

$$A(Q, Q', K) = e^{\frac{-(Q' - Q)}{K}}$$

for  $\mathcal{T} \subseteq \mathbb{R}$ . The process  $\mathcal{D}$  continues via simulated annealing if  $e$  is of the form

$$e(\langle x_u \rangle_{0 \leq u \leq t}, \tau) = x_{t+1} = \begin{cases} x'_t & \text{if } \tau(x'_t) \leq \tau(x_t) \text{ or } r \leq A(\tau(x_t), \tau(x'_t), \kappa(t)), \\ x_t & \text{otherwise,} \end{cases}$$

where  $x'_t \sim neighbors(x_t)$  and  $r \sim \mathbb{P}$  are drawn at random for each call to  $e$ .

**Definition 4** (population-based optimization). Let  $\mathcal{X}$  be a state space. Let  $\mathcal{T}$  be a target space with total order  $\leq$ . Let  $\tau : \mathcal{X} \rightarrow \mathcal{T}$  be a target function. A tuple  $\mathcal{E} = (\mathcal{X}, \mathcal{T}, \tau, E, \langle X_t \rangle_{0 \leq t \leq g})$  is a population-based optimization process iff  $X_t \in \wp^*(\mathcal{X})$  for all  $t$  and  $E : \langle \wp^*(\mathcal{X}) \rangle \times (\mathcal{X} \rightarrow \mathcal{T}) \rightarrow \wp^*(\mathcal{X})$  is a possibly randomized, non-deterministic, or further parametrized function so that the population-based optimization run is produced by calling  $E$  repeatedly, i.e.,  $X_{t+1} = E(\langle X_u \rangle_{0 \leq u \leq t}, \tau)$  where  $X_0$  is given externally or chosen randomly.

**Definition 5** (population-based optimization (alternate)). An optimization process  $\mathcal{E} = (\mathcal{X}, \mathcal{T}, \tau, E, \langle X_t \rangle_{0 \leq t \leq g})$  is called population-based iff  $\mathcal{X}$  has the form  $\mathcal{X} = \wp^*(\mathcal{Y})$  for some other state space  $\mathcal{Y}$ .

**Algorithm 4** (basic evolutionary algorithm). Let  $\mathcal{E} = (\mathcal{X}, \mathcal{T}, \tau, E, \langle X_u \rangle_{0 \leq u \leq t})$  be a population-based optimization process. The process  $\mathcal{E}$  continues via an evolutionary algorithm if  $E$  has the form

$$E(\langle X_u \rangle_{0 \leq u \leq t}, \tau) = X_{t+1} = selection(X_t \uplus variation(X_t))$$

where *selection* and *variation* are possibly randomized or non-deterministic functions so that for any  $X \in \wp^*(\mathcal{X})$  it holds that  $|selection(X)| \leq |X|$  and  $|selection(X \uplus variation(X))| = |X|$ .

**Theorem 1** (no free lunch [1, 2]). As measured by sample efficiency, i.e., the achieved minimal value of  $\tau$  per evaluations of  $\tau(x)$  for some new  $x \in \mathcal{X}$  for finite  $\mathcal{X}$ , all optimization algorithms perform the same when averaged over all possible target functions  $\tau$ . So, for any search/optimization algorithm, any elevated performance over one class of problems is exactly paid for in performance over another class.

**Definition 6** (multi-objective optimization). Let  $\mathcal{E} = (\mathcal{X}, \mathcal{T}, \tau, E, \langle X_u \rangle_{0 \leq u \leq t})$  be an optimization process.  $\mathcal{E}$  is a multi-objective optimization process iff the target space  $\mathcal{T}$  has the form  $\mathcal{T} = \mathcal{T}_0 \times \cdots \times \mathcal{T}_{N-1}$  for some  $N \in \mathbb{N}$  with  $<_i$  being a total order on  $\mathcal{T}_i$  for any  $i \in [0; N-1] \subset \mathbb{N}$ . Unless stated otherwise, we assume that no single total order on  $\mathcal{T}$  is available. However, we can construct a partial order  $\prec$  between target values  $g, g' \in \mathcal{T}$  so that

$$(g_0, \dots, g_{N-1}) \prec (g'_0, \dots, g'_{N-1}) \iff \forall i \in [0; N-1] \subset \mathbb{N} : g_i < g'_i,$$

which is sufficient to adapt many standard optimization algorithms.

**Definition 7** (Pareto front for optimization). Let  $\mathcal{E} = (\mathcal{X}, \mathcal{T}, \tau, E, \langle X_u \rangle_{0 \leq u \leq t})$  be a multi-objective optimization process with  $\prec$  being a partial order on the multi-objective target space  $\mathcal{T}$ .

- A solution candidate  $x$  Pareto-dominates a solution candidate  $x'$ , written  $x \prec x'$  (assuming minimization), iff  $\tau(x) \prec \tau(x')$ .
- A solution candidate  $x$  is Pareto-optimal if there exists no other solution candidate  $x' \in \mathcal{X}$  so that  $x' \prec x$ .
- The set of all Pareto-optimal solution candidates in  $\mathcal{X}$  is called the Pareto front of  $\mathcal{X}$  (w.r.t.  $\prec$ ).

**Algorithm 5** (gradient descent). Let  $\mathcal{E} = (\mathcal{X}, \mathcal{T}, \tau, e, \langle x_u \rangle_{0 \leq u \leq t})$  be an optimization process. Let  $\mathcal{T}$  be continuous ( $\mathcal{T} = \mathbb{R}$ , e.g.) and let  $\tau' : \mathcal{X} \rightarrow \mathcal{T}$  be the first derivative of  $\tau$ . The process  $\mathcal{E}$  continues via gradient descent (with learning rate  $\alpha \in \mathbb{R}^+$ ) if  $e$  is of the form

$$e(\langle x_u \rangle_{0 \leq u \leq t}, \tau) = x_{t+1} = x_t - \alpha \cdot \tau'(x_t).$$

The learning rate  $\alpha$  can also be given as a function, usually  $\alpha : \mathbb{N} \rightarrow \mathbb{R}$ , so that  $e(\langle x_u \rangle_{0 \leq u \leq t}, \tau) = x_{t+1} = x_t - \alpha(t) \cdot \tau'(x_t)$ .

If the computation of  $\tau'$  is stochastic, usually because it is derived from a stochastic sampling of data points, this process is called stochastic gradient descent (SGD).

**Algorithm 6** (gradient descent (policy)). Let  $\pi_\theta$  be a policy  $\pi$  that depends on a vector of continuous parameters  $\theta \in \Theta$ , usually with  $\Theta = \mathbb{R}^N$  for some  $N$ . Let  $\tau : \Theta \rightarrow \mathcal{T}$  be a target function on the parameters  $\theta$  of a policy  $\pi_\theta$ . Let  $\mathcal{T}$  be continuous ( $\mathcal{T} = \mathbb{R}$ , e.g.) and let  $\tau' : \Theta \rightarrow \mathcal{T}$  be the first derivative of  $\tau$ , i.e.,  $\tau'(\theta) = \frac{\partial \tau(\theta)}{\partial \theta}$ . If  $\mathcal{E} = (\Theta, \mathcal{T}, \tau, e, \langle x_u \rangle_{0 \leq u \leq t})$  is an optimization process that continues via gradient descent,  $\mathcal{E}$  is a process of policy optimization via gradient descent.

## References

- [1] No free lunch theorems. <http://www.no-free-lunch.org>. Accessed 2022-05-10.
- [2] David H Wolpert and William G Macready. No free lunch theorems for optimization. *IEEE transactions on evolutionary computation*, 1(1):67–82, 1997.