

# Computational Intelligence



LUDWIG-  
MAXIMILIANS-  
UNIVERSITÄT  
MÜNCHEN

LMU Munich  
winter term 2024/2025

Thomas Gabor  
Claudia Linnhoff-Popien

## schedule update

49	2024-12-03 Lecture #8	2024-12-05 Lecture #9
50	2024-12-10 Writing Exercise #4	2024-12-12 Writing Exercise #5
51	2024-12-17 Lecture #10	2024-12-19 Reading Exercise #3

recap

**Algorithm 5** (gradient descent). Let  $\mathcal{D} = (\mathcal{X}, \mathcal{T}, \tau, e, \langle x_u \rangle_{0 \leq u \leq t})$  be an optimization process. Let  $\mathcal{T}$  be continuous ( $\mathcal{T} = \mathbb{R}$ , e.g.) and let  $\tau' : \mathcal{X} \rightarrow \mathcal{T}$  be the first derivative of  $\tau$ . The process  $\mathcal{D}$  continues via gradient descent (with update rate  $\alpha \in \mathbb{R}^+$ ) if  $e$  is of the form

$$e(\langle x_u \rangle_{0 \leq u \leq t}, \tau) = x_{t+1} = x_t - \alpha \cdot \tau'(x_t).$$

The learning rate  $\alpha$  can also be given as a function, usually  $\alpha : \mathbb{N} \rightarrow \mathbb{R}$  so that  $e(\langle x_u \rangle_{0 \leq u \leq t}, \tau) = x_{t+1} = x_t - \alpha(t) \cdot \tau'(x_t)$ . If  $\tau$  is stochastic, this process is called stochastic gradient descent (SGD).

**Algorithm 6** (gradient descent (policy)). Let  $\pi_\theta$  be a policy  $\pi$  that depends on vector of continuous parameters  $\theta \in \Theta$  such that usually  $\Theta = \mathbb{R}^N$  for some  $N$ . Let  $\tau : \Theta \rightarrow \mathcal{T}$  be a target function on the parameters  $\theta$  of a policy  $\pi_\theta$ . Let  $\mathcal{T}$  be continuous ( $\mathcal{T} = \mathbb{R}$ , e.g.) and let  $\tau' : \Theta \rightarrow \mathcal{T}$  be the first derivative of  $\tau$ , i.e.,  $\tau'(\theta) = \frac{\partial \tau(\theta)}{\partial \theta}$ . If  $\mathcal{D} = (\Theta, \mathcal{T}, \tau, e, \langle x_u \rangle_{0 \leq u \leq t})$  is an optimization process that continues via gradient descent,  $\mathcal{D}$  is a process of policy optimization via gradient descent.

# Many Variants of Gradient Descent

recap

source:

<https://optimization.cbe.cornell.edu/index.php?title=AdaGrad>

---

**Algorithm 1:** AdaGrad general algorithm

---

```
 $\eta$ : Stepsize ;  
 $f(x)$ : Stochastic objective function ;  
 $x_1$ : Initial parameter vector;  
for  $t = 1$  to  $T$  do  
    Evaluate  $f_t(x_t)$  ;  
    Get and save  $g_t$  ;  
     $G_t \leftarrow \sum_{\tau=1}^t g_\tau g_\tau^\top$  ;  
     $x_{t+1} \leftarrow x_t - \eta G_t^{-1/2} g_t$  ;  
end  
return  $x_t$ 
```

---

source: <https://arxiv.org/pdf/1412.6980.pdf%5D>

---

**Algorithm 1:** *Adam*, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation.  $g_t^2$  indicates the elementwise square  $g_t \odot g_t$ . Good default settings for the tested machine learning problems are  $\alpha = 0.001$ ,  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$  and  $\epsilon = 10^{-8}$ . All operations on vectors are element-wise. With  $\beta_1^t$  and  $\beta_2^t$  we denote  $\beta_1$  and  $\beta_2$  to the power  $t$ .

---

**Require:**  $\alpha$ : Stepsize

**Require:**  $\beta_1, \beta_2 \in [0, 1]$ : Exponential decay rates for the moment estimates

**Require:**  $f(\theta)$ : Stochastic objective function with parameters  $\theta$

**Require:**  $\theta_0$ : Initial parameter vector

$m_0 \leftarrow 0$  (Initialize 1<sup>st</sup> moment vector)

$v_0 \leftarrow 0$  (Initialize 2<sup>nd</sup> moment vector)

$t \leftarrow 0$  (Initialize timestep)

**while**  $\theta_t$  not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$  (Get gradients w.r.t. stochastic objective at timestep  $t$ )

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$  (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$  (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$  (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$  (Compute bias-corrected second raw moment estimate)

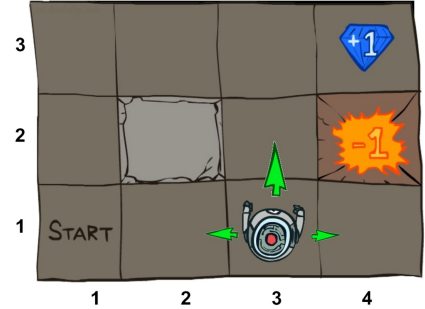
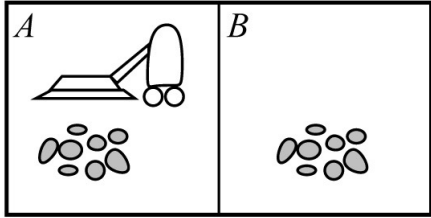
$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$  (Update parameters)

**end while**

**return**  $\theta_t$  (Resulting parameters)

---

## running examples



Russel, Norvig. Artificial Intelligence – A Modern Approach. Third Edition. 2016.  
[https://inst.eecs.berkeley.edu/~cs188/fa22/  
chatgpt.com](https://inst.eecs.berkeley.edu/~cs188/fa22/chatgpt.com)

Computational Intelligence, winter term 2024/2025, LMU Munich

encoding policies...

# Differentiable Programming



Jax (Python)

<https://github.com/google/jax>



Zygote (Julia)

```
julia> using Zygote  
  
julia> f(x) = 5x + 3  
  
julia> f(10), f'(10)  
(53, 5.0)
```

<https://github.com/FluxML/Zygote.jl>

# Neural Networks



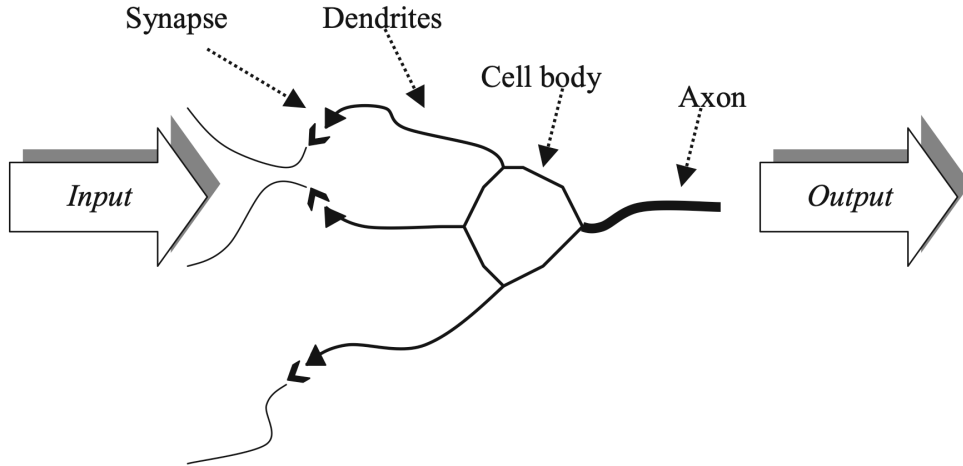
<https://pytorch.org>

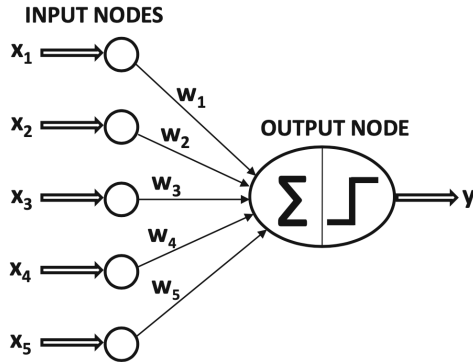


<https://www.tensorflow.org>

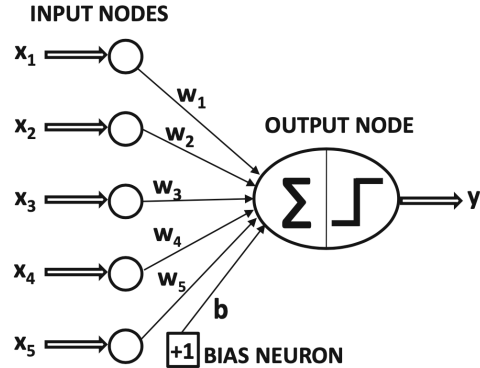


# Neural Networks

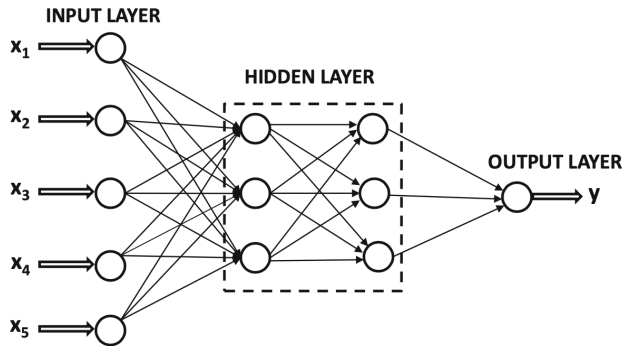




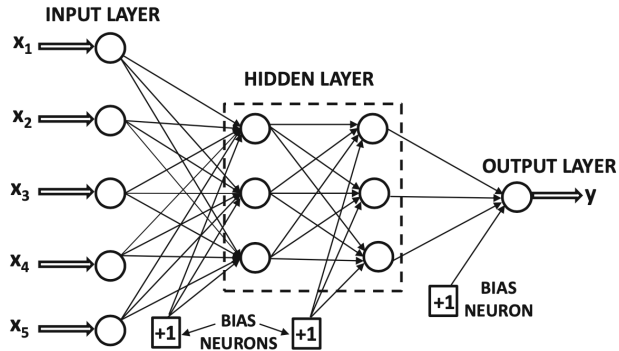
(a) Perceptron without bias



(b) Perceptron with bias

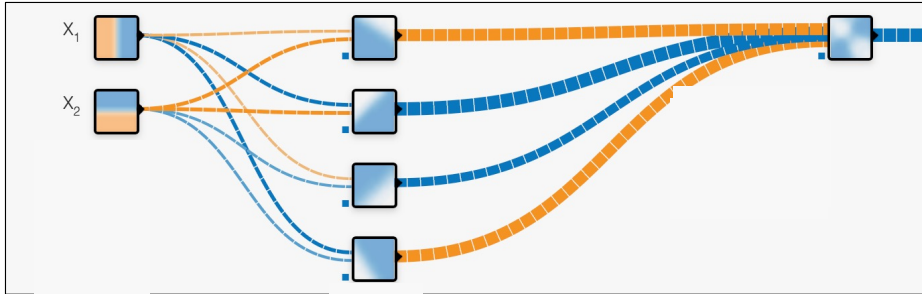


(a) No bias neurons



(b) With bias neurons

Let's try!



<https://playground.tensorflow.org>

**Definition 8** (neural network). A neural network (NN) is a function  $\mathcal{N} : \mathbb{R}^p \rightarrow \mathbb{R}^q$  with  $p$  inputs and  $q$  outputs. This function is defined via a graph made up of  $r$  layers  $L_1, \dots, L_r$  where each layer  $L_l$  consists of  $|L_l|$  cells  $C_{l,1}, \dots, C_{l,|L_l|}$ , which make up the graph's vertices, and each cell  $C_{l,c}$  of the layer  $L_l$  is connected to all cells of the previous layer, i.e.,  $C_{l-1,d}$  for  $d = 1, \dots, |L_{l-1}|$ , via the graph's edges. Each edge of a cell  $C_{l,c}$  is assigned an edge weight  $E_{l,c,e} \in \mathbb{R}, e = 1, \dots, |L_{l-1}|$ . Given a fixed graph structure and activation function  $f : \mathbb{R} \rightarrow \mathbb{R}$ , the vector of all edge weights

$$\mathbf{w} = \langle E_{l,c,e} \rangle_{l=1,\dots,r, \ c=1,\dots,|L_l|, \ e=1,\dots,|L_{l-1}|}$$

and the vector of all cell biases

$$\mathbf{b} = \langle B_{l,c} \rangle_{l=1,\dots,r, \ c=1,\dots,|L_l|}$$

with  $B_{l,c} \in \mathbb{R}$  define the network's functionality. The combined vector  $\overline{\mathcal{N}} = \mathbf{w} \# \mathbf{b}$  is called the network  $\mathcal{N}$ 's parameters.

**Definition 8** (neural network). A neural network (NN) is a function  $\mathcal{N} : \mathbb{R}^p \rightarrow \mathbb{R}^q$  with  $p$  inputs and  $q$  outputs. This function is defined via a graph made up of  $r$  layers  $L_1, \dots, L_r$  where each layer  $L_l$  consists of  $|L_l|$  cells  $C_{l,1}, \dots, C_{l,|L_l|}$ , which make up the graph's vertices, and each cell  $C_{l,c}$  of the layer  $L_l$  is connected to all cells of the previous layer, i.e.,  $C_{l-1,d}$  for  $d = 1, \dots, |L_{l-1}|$ , via the graph's edges. Each edge of a cell  $C_{l,c}$  is assigned an edge weight  $E_{l,c,e} \in \mathbb{R}, e = 1, \dots, |L_{l-1}|$ . Given a fixed graph structure and activation function  $f : \mathbb{R} \rightarrow \mathbb{R}$ , the vector of all edge weights

$$\mathbf{w} = \langle E_{l,c,e} \rangle_{l=1,\dots,r, \ c=1,\dots,|L_l|, \ e=1,\dots,|L_{l-1}|}$$

and the vector of all cell biases

$$\mathbf{b} = \langle B_{l,c} \rangle_{l=1,\dots,r, \ c=1,\dots,|L_l|}$$

with  $B_{l,c} \in \mathbb{R}$  define the network's functionality. The combined vector  $\overline{\mathcal{N}} = \mathbf{w} \# \mathbf{b}$  is called the network  $\mathcal{N}$ 's parameters.

A network's output given an input  $\mathbf{x} \in \mathbb{R}^p$  is given via

$$\mathbf{y} = \mathcal{N}(\mathbf{x}) = \langle O(r, c) \rangle_{c=1,\dots,|L_r|} \in \mathbb{R}^q$$

$$\text{where } O(l, c) = \begin{cases} x_c & \text{if } l = 0, \\ f(B_{l,c} + \sum_{i=1}^{|L_{l-1}|} E_{l,c,i} \cdot O(l-1, i)) & \text{otherwise.} \end{cases}$$