

Computational Intelligence



LUDWIG-
MAXIMILIANS-
UNIVERSITÄT
MÜNCHEN

LMU Munich
winter term 2024/2025


Thomas Gabor
Claudia Linnhoff-Popien

First: What are states?

Second: What are state values?
(encoding policies)

running example #1

The Vacuum World




Reward function: +1 if dirty room
becomes clean

\Rightarrow states after we cleaned often
are good

running example #2

The Basic Grid World



reward function: get reward according to grid field

<https://inst.eecs.berkeley.edu/~cs188/fa22/>

Computational Intelligence, winter term 2024/2025, LMU Munich

233

running example #3

Resource/Stock Trading



running example #4

Personal Life Assistant



Third: Where do we get
a state value function?

(finding policies)

Theorem 2 (Bellman equation). Let $(\mathcal{S}, \mathcal{A}, \mathcal{T}, P, R)$ be a Markov decision process. Let $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{T}$ be the expected reward of executing an action in a given state, i.e., $R(s, a) = \mathbb{E}[R(s, a, s')]$ where $s' \sim P(s'|s, a)$. Let $\gamma \in [0; 1) \subseteq \mathbb{R}$ be a temporal discount factor.

The expected reward of a policy π being executed starting from state s is given via π 's value function

$$V^\pi(s) = R(s, \pi(s)) + \gamma \cdot \sum_{s' \in \mathcal{S}} P(s'|s, \pi(s)) \cdot V^\pi(s').$$

The value function of the optimal policy π^* is given via

$$V^{\pi^*}(s) = \max_{a \in \mathcal{A}} \left(R(s, a) + \gamma \cdot \sum_{s' \in \mathcal{S}} P(s'|s, a) \cdot V^{\pi^*}(s') \right).$$

Theorem 2 (Bellman equation). Let $(\mathcal{S}, \mathcal{A}, \mathcal{T}, P, R)$ be a Markov decision process. Let $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{T}$ be the expected reward of executing an action in a given state, i.e., $R(s, a) = \mathbb{E}[R(s, a, s')]$ where $s' \sim P(s'|s, a)$. Let $\gamma \in [0; 1) \subseteq \mathbb{R}$ be a temporal discount factor.

The expected reward of a policy π being executed starting from state s is given via π 's value function

$$V^\pi(s) = R(s, \pi(s)) + \gamma \cdot \sum_{s' \in \mathcal{S}} P(s'|s, \pi(s)) \cdot V^\pi(s').$$

The value function of the optimal policy π^* is given via

$$V^{\pi^*}(s) = \max_{a \in \mathcal{A}} \left(R(s, a) + \gamma \cdot \sum_{s' \in \mathcal{S}} P(s'|s, a) \cdot V^{\pi^*}(s') \right).$$

Algorithm 7 (optimal policy). Let $V^* : \mathcal{S} \rightarrow \mathcal{T}$ be the *true value function* of a Markov decision process $(\mathcal{S}, \mathcal{A}, \mathcal{T}, P, R)$. The optimal policy $\pi^* : \mathcal{S} \rightarrow \mathcal{A}$ is given via


$$\pi^*(s_t) = \arg \max_{a \in \mathcal{A}} V^*(s')$$

where $s' \sim P(s'|s_t, a)$ is the follow-up state when executing action a in state s_t .

Let's try that!


running example #1

The Vacuum World



running example #2

The Basic Grid World



running example #3

Resource/Stock Trading



running example #4

Personal Life Assistant



Putting it together...

Definition 11 (training of a neural network). Let $\mathcal{N} : \mathbb{R}^p \rightarrow \mathbb{R}^q$ be a neural network with n weights $\overline{\mathcal{N}} = \mathbf{w} \# \mathbf{b} \in \mathbb{R}^n$ as in Definition 8. Note that thus $|\overline{\mathcal{N}}| = n$. Let $\tau : \mathbb{R}^n \rightarrow \mathbb{R}$ be a target function as in Definition 2. Note that thus $\mathcal{T} = \mathbb{R}$. The process of optimizing the network weights $\overline{\mathcal{N}}$ so that $\tau(\overline{\mathcal{N}})$ becomes minimal is called training.

- Let $\mathbb{T} = \{(\mathbf{x}_i, \mathbf{y}_i) : i = 1, \dots, N\}$ be a set of N points of training data, where $\mathbf{x}_i \in \mathbb{R}^p, \mathbf{y}_i \in \mathbb{R}^q$ for all i .
If τ is of the form

$$\tau(\overline{\mathcal{N}}) = \sum_{i=1}^N (\mathcal{N}(\mathbf{x}_i) - \mathbf{y}_i)^2$$

or a similar form, the process of training \mathcal{N} is called supervised learning.

Definition 11 (training of a neural network). Let $\mathcal{N} : \mathbb{R}^p \rightarrow \mathbb{R}^q$ be a neural network with n weights $\overline{\mathcal{N}} = \mathbf{w} \# \mathbf{b} \in \mathbb{R}^n$ as in Definition 8. Note that thus $|\overline{\mathcal{N}}| = n$. Let $\tau : \mathbb{R}^n \rightarrow \mathbb{R}$ be a target function as in Definition 2. Note that thus $\mathcal{T} = \mathbb{R}$. The process of optimizing the network weights $\overline{\mathcal{N}}$ so that $\tau(\overline{\mathcal{N}})$ becomes minimal is called training.

- Let $\mathbb{T} = \{(\mathbf{x}_i, \mathbf{y}_i) : i = 1, \dots, N\}$ be a set of N points of training data, where $\mathbf{x}_i \in \mathbb{R}^p, \mathbf{y}_i \in \mathbb{R}^q$ for all i .
If τ is of the form

$$\tau(\overline{\mathcal{N}}) = \sum_{i=1}^N (\mathcal{N}(\mathbf{x}_i) - \mathbf{y}_i)^2$$

or a similar form, the process of training \mathcal{N} is called supervised learning.

- Let $(\mathcal{O}, \mathcal{A}, \mathcal{T}, e, R)$ be a decision process (cf. Definition 9) for which policy $\pi_{\overline{\mathcal{N}}} : \mathcal{O} \rightarrow \mathcal{A}$ yields (possibly randomized or non-deterministic) rewards $\langle r_t \rangle_{t \in \mathbb{Z}}$. Note that $\pi_{\overline{\mathcal{N}}}$ in some way calls \mathcal{N} to produce its output, for example

$$\pi_{\overline{\mathcal{N}}}(o) = \mathcal{N}(o)$$

for $\mathcal{O} \subseteq \mathbb{R}^p, \mathcal{A} \subseteq \mathbb{R}^q$ or if suitable translations exist.

If τ is of the form

$$\tau(\overline{\mathcal{N}}) = -\mathbb{E} \left[\sum_{t \in \mathbb{Z}} \gamma^t \cdot r_t \right]$$

or a similar form, the process of training \mathcal{N} is called policy-based reinforcement learning.

- Let $(\mathcal{S}, \mathcal{A}, \mathcal{T}, P, R)$ be a Markov decision process (cf. Definition 10) for which we run policy $\pi_{\overline{\mathcal{N}}} : \mathcal{S} \rightarrow \mathcal{A}$. Note that $\pi_{\overline{\mathcal{N}}}$ in some way calls \mathcal{N} to produce its output, for example

$$\pi_{\overline{\mathcal{N}}}(s) = \arg \max_{a \in \mathcal{A}} \mathbb{E}_{s' \sim P(s'|s, a)} [\mathcal{N}(s')]$$

for $\mathcal{S} \times \mathcal{A} \subseteq \mathbb{R}^p$ with $q = 1$ or if suitable translations exist.

Let $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{T}$ be the expected reward of executing an action in a given state, i.e., $R(s, a) = \mathbb{E}[R(s, a, s')]$ where $s' \sim P(s'|s, a)$. Let $T : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$ be a (possibly randomized or non-deterministic) transition function, i.e., $T(s, a) = s'$ where $s' \sim P(s'|s, a)$. Let $\gamma \in [0; 1]$ be a discount factor. Let $V_{\pi_{\overline{\mathcal{N}}}} : \mathcal{S} \rightarrow \mathbb{R}$ be the total discounted reward that policy $\pi_{\overline{\mathcal{N}}}$ generates when starting in state s , i.e.,

$$V_{\pi_{\overline{\mathcal{N}}}}(s) = R(s, \pi_{\overline{\mathcal{N}}}(s)) + \gamma \cdot V_{\pi_{\overline{\mathcal{N}}}}(T(s, \pi_{\overline{\mathcal{N}}}(s))).$$

Note that for $\gamma < 1$ we can abort this recursive computation once the effect of the further recursive part is sufficiently small. Note that we may also have a fixed recursion depth or that $T(s^\dagger, \cdot)$ might not be defined for all $s^\dagger \in \mathcal{S}$, which are then called terminal states and also cause the recursion to end.

Let $\mathbb{S} = \{\mathbf{s}_i : i = 1, \dots, N\} \subseteq \mathcal{S}$ be a set of training states. If τ is of the form

$$\tau(\overline{\mathcal{N}}) = -\frac{1}{N} \cdot \sum_{i=1}^N V_{\pi_{\overline{\mathcal{N}}}}(\mathbf{s}_i)$$

or a similar form, the process of training \mathcal{N} is called value-based reinforcement learning.

Reinforcement Learning

Variations of Value Functions

name	network	policy
policy-based	$\mathcal{N} : \mathcal{S} \rightarrow \mathcal{A}$	$\pi_{\overline{\mathcal{N}}}(s) = \mathcal{N}(s)$
value-based (V)	$\mathcal{N} : \mathcal{S} \rightarrow \mathbb{R}$	$\pi_{\overline{\mathcal{N}}}(s) = \arg \max_{a \in \mathcal{A}} \mathbb{E}_{s' \sim P(s' s,a)} [\mathcal{N}(s')]$
value-based (Q)	$\mathcal{N} : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$	$\pi_{\overline{\mathcal{N}}}(s) = \arg \max_{a \in \mathcal{A}} \mathcal{N}(s, a)$

...

On-Policy vs. Off-Policy Learning

The Goal Class Hierarchy

